

# COMP6771

## Advanced C++ Programming

### Week 5

### Part Two: Dynamic Memory Management

2016

[www.cse.unsw.edu.au/~cs6771](http://www.cse.unsw.edu.au/~cs6771)

# Initialisation Revisited

```
1  #include <iostream>
2
3  struct X {
4      X() { std::cout << "X() "; }
5  };
6
7  struct Y {
8      Y() { std::cout << "Y() "; }
9  };
10
11 class A {
12     X x;
13 public:
14     A() { std::cout << "A() "; }
15 };
16
17 class B : public A {
18     Y y;
19 public:
20     B() { std::cout << "B() "; }
21 };
22
23 int main() {
24     B b;
25 }
```

What is the output?

- a. X() A() Y() B()
- b. A() X() B() Y()
- c. Y() B() X() A()
- d. B() Y() A() X()
- e. X() Y() A() B()

# Initialisation Revisited

```
1  #include <iostream>
2
3  struct X {
4      X() { std::cout << "X() "; }
5  };
6
7  struct Y {
8      Y() { std::cout << "Y() "; }
9  };
10
11 class A {
12     X x;
13 public:
14     A() { std::cout << "A() "; }
15 };
16
17 class B : public A {
18     Y y;
19 public:
20     B() { std::cout << "B() "; }
21 };
22
23 int main() {
24     B b;
25 }
```

What is the output?

a. X() A() Y() B()

b. A() X() B() Y()

c. Y() B() X() A()

d. B() Y() A() X()

e. X() Y() A() B()

# Finalisation Revisited

```
1  #include <iostream>
2
3  struct X {
4      ~X() { std::cout << "~X() "; }
5  };
6
7  struct Y {
8      ~Y() { std::cout << "~Y() "; }
9  };
10
11 class A {
12     X x;
13 public:
14     ~A() { std::cout << "~A() "; }
15 };
16
17 class B : public A {
18     Y y;
19 public:
20     ~B() { std::cout << "~B() "; }
21 };
22
23 int main() {
24     B b;
25 }
```

What is the output?

# Finalisation Revisited

```
1  #include <iostream>
2
3  struct X {
4      ~X() { std::cout << "~X() "; }
5  };
6
7  struct Y {
8      ~Y() { std::cout << "~Y() "; }
9  };
10
11 class A {
12     X x;
13 public:
14     ~A() { std::cout << "~A() "; }
15 };
16
17 class B : public A {
18     Y y;
19 public:
20     ~B() { std::cout << "~B() "; }
21 };
22
23 int main() {
24     B b;
25 }
```

What is the output?

Reversed:

**~B() ~Y() ~A() ~X()**

# Lifetimes and Named/Unnamed Objects

- **Lifetime:** the period of time in which memory is allocated for an object
- Different kinds of objects:
  - Static objects: allocated in a global (static) area
  - Local (or automatic or stack) objects
  - Heap objects
- **Named objects:** (named by the programmer): their lifetimes determined by their scope
- **Heap objects (unnamed objects):** their lifetimes determined by the programmer

## Local Objects

```
void f() {  
    int i = 1;           // automatic  
}
```

- **Lifetime:** created when `f` is called and cease to exist when **this invocation** of `f` completes (**with destructors called**)
- Primitive objects are **not** initialised by default
- Class objects are initialised by constructors
- The same name denotes different objects during different invocations of the function `f`
- Memory management done by the compiler

## Local Objects (Cont'd)

```
1 {  
2     std::stack<int> s1;  
3     {  
4         std::stack<int> s2;  
5         } ---> call ~stack() for s2  
6     } ---> call ~stack() for s1
```

- Every object inaccessible at its closing '}'s
- The destructors s1 and s2 are called when s1 and s2 die

```
1 {  
2     std::stack<int> s1;  
3     {  
4         std::stack<int> &s2 = s1;  
5         } ---> ???  
6     } ---> call ~stack() for s1
```

- s2 is a reference to s1
- The destructor for s1 is called when s1 dies



## Temporary Objects

- Unnamed objects created on the stack by the compiler.
- Read the detailed rules: [http://www-01.ibm.com/support/knowledgecenter/SSGH3R\\_8.0.0/com.ibm.xlcpp8a.doc/language/ref/cplr382.htm%23cplr382](http://www-01.ibm.com/support/knowledgecenter/SSGH3R_8.0.0/com.ibm.xlcpp8a.doc/language/ref/cplr382.htm%23cplr382)
- Used during reference initialization and during evaluation of expressions including type conversions, argument passing, function returns, and evaluation of the throw expression.
- There are two exceptions in the destruction of full-expressions:
  - The expression appears as an initializer for a declaration defining an object: the temporary object is destroyed when the initialization is complete.
  - A reference is bound to a temporary object: the temporary object is destroyed at the end of the reference's lifetime.

## Lifetimes of Temporary Objects

```
1 #include<vector>
2
3 std::vector<int> f() { return std::vector<int>{}; }
4
5 void g() {
6     const std::vector<int>& v = f();
7 }
```

The lifetime of the temporary vector created f may persist as long as the reference is live.

# Static Objects

- Four categories of objects:
  - Local static objects
  - **Global objects**
  - Namespace objects
  - Class static objects
- Memory management done by the compiler

These objects are often referred to as global/static objects.

## (Local) Static Objects

```
int f() {  
    static int i = 0;  
    return ++i;  
}  
cout << f(); // prints 1  
cout << f(); // prints 2
```

- **lifetime**: created when `f` is **first** called and persist until the program completes (**with destructors called**)
- Primitive objects are initialised to default values
- Class objects are initialised by constructors
- (Local) static objects are still visible locally

## Global, Namespace and Class Static Objects


file.cpp:

```
std::stack<int> s;           // global
namespace X {
    std::queue<short> q;     // namespace
}
class Y {
    static std::list<char> l; // class static
}
```

- All created at program start-up and persist until the program completes (**with their destructors called**)
- Primitive objects initialised to default values
- Class objects are initialised by constructors

## Heap (or Free-Store) Objects

- Allocated via **new** and persist until deleted via **delete**
- In “`int *pi = new int{1}`”, **two** objects are involved:



- pi is destroyed at the end of its lifetime automatically
    - the pointed-to object must be destroyed explicitly
  - Thus, **important to remember**:
    - `delete pi`  $\equiv$  `delete *pi`  $\equiv$  `unnamed`
    - `delete pi`  $\neq$  `delete named`
- First, the destructor for the pointed-to object, `*p`, is called
  - Then, the corresponding memory is freed

The destructor for a built-in type does nothing.

## Heap Objects (Cont'd)

```
int* f() {  
    int *pi = new int{1};  
    return pi;  
}  
int main() {  
    int *pj = f();  
    delete pj;  
}
```

- pi and pj are named local objects
- new int(1) is an unnamed heap object

## Match delete with the Corresponding new

```
int *pi = new int{1};
int *pj = new int[10];
           // an array of 10 ints uninitialised
...

delete pi;    // delete [] pi is undefined
delete [] pj; // delete pj is undefined
```

An array object typically contains the information regarding the number of its elements



## new: Dynamically Allocate and Initialise Objects

```

1  int *pi = new int;           // an uninitialised int
2  int *pj = new int{};        // an int initialised to 0
3
4  std::string *ps1 = new std::string; // initialised to empty
5                                     // string by calling string's default ctor
6  std::string *ps2 = new std::string{"abc"}; // initialised to
7                                     // string "abc"
8
9  auto *pf = new auto{3.14};
10         // inferred as float *pf = new float{3.14};
11
12  const int *pci = new const int{100};
13         // an const int object initialised to 100
14
15  std::vector<int> *pv1 = new std::vector<int>{};
16         // a vector<int> initialised to be empty
17         // if allocation fails, throws std::bad_alloc
18
19  std::vector<int> *pv2 = new (nothrow) std::vector<int>{};
20         // a vector<int> initialised to be empty
21         // if allocation fails, returns nullptr

```

# Many Memory-Corruption Bugs in C/C++ Code!

```
1 int i, *pi1 = &i, *pi2 = nullptr;
2 double *pd = new double{3.14}, *pd2 = pd;
3
4 delete i;    // error: i is not a pointer
5 delete pi1; // undefined: pi1 points to a stack location
6 delete pd;  // ok
7 delete pd2; // double free!
8 *pd = 1.0;  // accessing a dangling pointer
9 delete pi2; // ok: deleting nullptr does nothing
10
11 int* foo() {
12     int *pi = new std::vector<int>{50};
13     return pi;
14 }
15
16 int *pj = foo();
17 // If the programmer never calls delete pj, then you have
18 // a memory leak! The memory storing the vector has leaked.
```

# Resetting the Value of a Pointer After a delete

- A deleted pointer is not automatically nullified
- Set a deleted pointer to **nullptr** yourself:

```
1 int *p = new int{1};  
2 delete p;  
3 p = nullptr;
```

- We nullify the pointers in a moved-from object:

```
1 UB_stack::UB_stack(const UB_stack &&s) {  
2     head_ = s.head_;  
3     s.head = nullptr; // put the moved-from object  
4 }                     // in a valid state to be destroyed
```

## new: Dynamic Arrays

- Allocate and initialise dynamic arrays

```
1 int *p = new int[5];    // block of 5 uninitialised ints
2 int *p = new int[5]{}; // block of 5 ints initialised to 0
3
4 // block of 5 empty strings
5 std::string *p = new std::string[5];
6 // block of 5 empty strings
7 std::string *p = new std::string[5]{};
8
9 int *p = new int[5](1);           // error
10 std::string *p = new std::string[5]("abc"); // error
11
12 int *p = new int[5]{0, 2, 3, 4, 5}; // ok
13 std::string *p = new std::string[5]{"an"};
14 // ok: the first initialised
15 // to "an" and the last four to empty strings
```

# Motivations for Memory Management

- Standard solution (value semantics)

```
SMatrix operator+(const SMatrix &a, const SMatrix &b) {
    SMatrix c = a;
    return c+=b;
} // can be efficient now in C++11 due to move ctor/assignment
```

- Efficient

```
SMatrix& operator+(const SMatrix &a, const SMatrix &b) {
    SMatrix c = a;
    return c+=b;
} // but incorrect
```

- Efficient

```
SMatrix& operator+(const SMatrix &a, const SMatrix &b) {
    SMatrix *c = new SMatrix;
    return *c+=b;
} // but who is to free the heap-allocated object?
```

- When objects are “large”, heap objects may be used. Using pointers is a burden on users  $\implies$  memory management techniques (smart pointers, reference counting, handles, etc.) so that users don’t have to manage pointers explicitly themselves

# Smart Pointers: Motivations

- Memory leak:

```
void f() {  
    X* x = new X{};  
    ...  
    delete x;  
}
```

- if x is not **deleted**, you suffer the **memory leak** problem
- If an exception is thrown in **...**, **delete x** will not be executed. You will then suffer the **memory leak** problem
- Double free and dangling pointers

## One Solution: Bad Programming Style

```
void f() {  
    X* x = new X{};  
    try {  
        ...  
    }  
    catch (...) {  
        delete x;  
        throw;        // rethrow the exception  
    }  
  
    delete x;  
}
```

## Smart Pointers: Overloading of $\rightarrow$

- Usual notations:
  - if  $p$  is a built-in pointer,  $p \rightarrow x$  accesses the member  $x$
  - if  $p$  is an object or a reference to an object,  $p.x$  accesses the member  $x$
- **Smart pointers:**
  - a class type exhibits a **pointer-like** behaviour when the  $\rightarrow$  is overloaded
  - If  $p$  is an object or reference,  $p \rightarrow x$  can be used
  - $p$  is known as a **smart pointer**
- Item 28, Meyers' More Effective C++



## Basic Idea: Defining a Smart Pointer to Strings

```
1 #include <iostream>
2
3 class StringPtr {
4 public:
5     StringPtr(std::string *p) : ptr{p} { }
6     ~StringPtr() { delete ptr; }
7     std::string* operator->() { return ptr; }
8     std::string& operator*() { return *ptr; }
9 private:
10     std::string *ptr;
11 };
12
13 int main() {
14     std::string *ps = new std::string{"smart pointer"};
15     StringPtr p{ps};
16     std::cout << *p << std::endl;
17     std::cout << p->size() << std::endl;
18 }
```

# Using (Named) Objects to Prevent Memory Leaks

- Place a pointer inside a named object and have the object manage the pointer for you

```
class StringPtr {
    ~StringPtr() { delete ptr; }
    ...
    std::string *ptr;
};
```

- User code:

```
{
StringPtr p(new std::string("smart pointer"));
} <-- *ptr freed here when ~StringPtr() called
```

- The object p's destructor will delete the heap string object

## What About Double Free?

- Yes, it is still with us:

```
class StringPtr {
    ~StringPtr() { delete ptr; }
    ...
    std::string *ptr;
};
```

- User code:

```
{
    StringPtr p(new std::string("smart pointer"));
    {
        StringPtr q(p);
    } <-- *ptr freed here
    std::cout << p->size() // Oops
} <-- *ptr freed again
```

## Reference Counting (RC)

- A classic automatic garbage collection technique
- Can be used manually by the C++ programmer to prevent
  - memory leaks, and
  - the dangling pointer problem
- FAQs 31.10 – 31.11: RC with copy-on-write semantics
- Item 29, Meyers' More Effective C++

## A Simple RC Scheme (without COW)

```
1  class PointPtr;  
2  
3  class Point {  
4  public:  
5      friend class PointPtr;  
6      static PointPtr create();  
7      static PointPtr create(int i, int j);  
8  private:  
9      Point() : count_(0) { i = j = 0; }  
10     Point(int i, int j) : count_(0) {  
11         this->i = i; this->j = j;  
12     }  
13     unsigned count_;  
14     int i, j;  
15 };
```

- Users cannot construct point objects directly
- Factory methods used to construct point objects

## A Simple RC Scheme (Cont'd)

```

1 class PointPtr { // smart pointer
2 public:
3     Point* operator-> () { return p_; }
4     Point& operator* () { return *p_; }
5     PointPtr(Point* p) : p_{p} { ++p_->count_; } // p cannot be NULL
6     ~PointPtr() { if (--p_->count_ == 0) delete p_; }
7     PointPtr(const PointPtr& p) : p_{p.p_} { ++p_->count_; }
8     PointPtr& operator= (const PointPtr& p) {
9         ++p.p_->count_;
10        if (--p_->count_ == 0) delete p_;
11        p_ = p.p_;
12        return *this;
13    }
14 private:
15     Point* p_; // p_ is never NULL
16 };
17 inline PointPtr Point::create() { return new Point(); }
18 inline PointPtr Point::create(int i, int j) {
19     return new Point(i, j);
20 }

```

## A Simple RC Scheme (Cont'd)

- Client code:

```
int main() {  
    PointPtr p1 = Point::create(1, 1);  
    PointPtr p2 = p1;  
    p1 = Point::create(2, 2);  
    p2 = p1;  
}
```

- In the two overloaded `Point::create`, the constructor `PointPtr (Point *p)` is called on return

# Does new Do More Than Allocate Memory (FAQ12.01)?

- Consider:

```
1 X *p = new X();
```

- The compiler-generated code looks something like:

```
1 X *p = (X*) operator new( sizeof(X) ); //Step 1
2 try {
3     new (p) X(); // Step 2
4 } catch (...) {
5     operator delete(p);
6     throw;
7 }
```

- No memory leak occurs if the constructor throws an exception



# Does delete Do More Than Deallocate Memory (FAQ12.01)?

- Consider:

```
1 delete p // p is an object of class X
```

- The compiler-generated code looks something like:

```
1 p->~X(); //Step 1: calling the destructor
2 operator delete(p); //Step 2: freeing the memory
```

- Deleting a null pointer is safe


```
1 if (p != 0) delete p; // test redundant
2 // the test done in operator delete(void *p)
```

## A Variety of Smarter Pointers!

- **Smart pointers:** Strategies for handling  $p = q$ , where  $p$  and  $q$  are smarter pointers
- “New” smart pointers introduced in C++11:
  - `unique_ptr`
  - `shared_ptr`
  - `weak_ptr`
- Why smarter pointers?
  - Less bugs (e.g., fewer double-free & dangling pointer problems)
  - Exception safety (e.g., automatic memory deallocation)
  - Gabagge collection (RC)
  - Efficiency (copy-on-write)

## unique\_ptr (§12.1.5)

- A `unique_ptr` **owns** the object to which it points to
- The underlying object always owned by one owner only
- When the `unique_ptr` goes out of scope, its internal object pointed to by its **internal pointer** is destroyed



- Explicit ownership transfer:

```

1 std::unique_ptr<UB_stack> p(new UB_stack{});
2 std::unique_ptr<UB_stack> q1 = p; // error: no copy
3 std::unique_ptr<UB_stack> q2;
4 q2 = p; // error: no assignment
5 std::unique_ptr<UB_stack> q3;
6 q3.reset(p.release()); // ok: ownership transfer
  
```

## Passing and Returning `unique_ptr`

```
1 std::unique_ptr<UB_stack> foo() {  
2     std::unique_ptr<UB_stack> p(new UB_stack{});  
3     p.get().push(1);  
4     p.get().push(2);  
5     return p;  
6 }  
7  
8 int main() {  
9     std::unique_ptr<UB_stack> q = foo();  
10 }
```

- Can assign/copy a temporary `unique_ptr` that is about to die
- The compiler makes sure that `q` will be the unique owner after the call to `foo` returned; the temporary has lost the ownership

# Smarter Pointers and Exception

- Smarter pointers:

```
1 void f() {  
2     std::unique_ptr<UB_stack> up(new UB_stack{});;  
3  
4     code that throws an exception  
5  
6 } // the destructor for up called here  
7     // the underlying stack object is freed
```

- Dumb pointers:

```
1 void f() {  
2     UB_stack *p = new UB_stack();  
3  
4     code that throws an exception  
5  
6     delete p;  
7 } // the destructor for *p not called here  
8     // the underlying stack object has leaked
```

The destructors for all named objects are called only!

## unique\_ptr

```

1 std::unique_ptr<int> x = new int{i};
2 std::unique_ptr<int> y = x; // Compile error
3 std::unique_ptr<int> z = std::move(x); // Transfers ownership
4                                     // z now owns the memory
5                                     // x is rendered invalid
6 z.reset(); // Deletes the memory

```

- copy constructor and assignment disabled
- **move** used to transfer ownership
- Create more than one owner (**Don't do this!**):

```

{
    int *pi = new int{1};
    std::unique_ptr<int> p ( pi );
    std::unique_ptr<int> q ( pi );
}    --> new int(1) will be deleted twice

```

Two unique\_ptr objects should not hold ownership to the same heap object. **The double-free problem!**



## shared\_ptr

```
std::shared_ptr<int> x = std::make_shared<int>(1);
std::shared_ptr<int> y = x; // Both now own the memory
```

```
x.reset(); // Memory still exists, due to y.
y.reset(); // Deletes the memory, since
            // no one else owns the memory
```

- Reference counted ownership: every copy of the same **shared\_ptr** owns the same pointed-to object
- Freed only when all instances are destroyed
- What about cycles?  
(e.g., A has shared\_ptr to B, B has shared\_ptr to A)



## weak\_ptr

```

1  std::shared_ptr<int> x = std::make_shared<int>(1);
2  std::weak_ptr<int> wp = x; // x owns the memory
3  {
4      std::shared_ptr<int> y = wp.lock(); // x and y own the memory
5      if (y)
6          { // Do something with y }
7  } // y is destroyed. Memory is owned by x
8
9  x.reset(); // Memory is deleted
10
11 std::shared_ptr<int> z = wp.lock(); // Memory gone; get nullptr
12 if (z)
13     { // will not execute this }

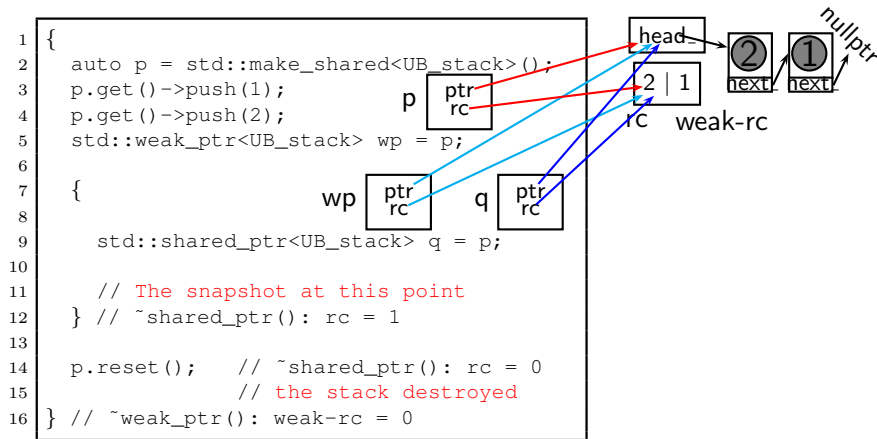
```

- Useful when you need to access to a resource through a pointer that can outlive the resource
- Read:

<http://www.drdobbs.com/weak-pointers/184402026>

## weak\_ptr + shared\_ptr (§12.1.6)

- Points to an object managed by a shared\_ptr
- Doesn't get reference-counted



- Conceptually illustrated but may vary across implementations

## Smart Pointers for Arrays

- Different operations provided for arrays (§12.2.1)

```
1 #include<iostream>
2 #include<memory>
3 #include "UB_stack.h"
4
5 int main() {
6     auto p =std::unique_ptr<UB_stack[]>(5);
7     p[0].push(1);
8     p[1].push(2);
9 } // <-- delete[] called on the underlying array object
```

- `shared_ptr` does not provide support for arrays. You can supply your own deleter (Page 480) - **Do not do this!**:

```
1 struct deleter {
2     void operator ()( UB_stack const * p) { delete[] p; }
3 };
4
5 int main() {
6     std::shared_ptr<UB_stack> p (new UB_stack[5](), deleter());
7 }
```

# Reading

- Section 18.1
- C++ Memory and Resource Management  
<http://www.informit.com/articles/article.asp?p=30642>
- C++ FAQ Lite: <http://yosefk.com/c++fqa/exceptions.html>
- How is exception handling implemented?  
<http://www.codeproject.com/cpp/exceptionhandler.asp>