

COMP6771

Advanced C++ Programming

Week 9

Multithreading (continued)

2016

www.cse.unsw.edu.au/~cs6771

So Far

- **Program Workflows:** Sequential, Parallel, Embarrassingly Parallel
- **Memory:** Shared Memory, Local Memory, Message Passing
- **Atomic Operations:** Generally one CPU clock cycle
- **Concepts and Problems:** Race Conditions, Mutual Exclusion, Mutex Objects and Deadlocks

threadLambdaRace.cpp

```
1  #include <iostream>
2  #include <thread>
3
4  // NOTE: do not compile this with -O2 it optimises out the ++
5  // call and prevents the race condition from happening.
6
7  int main() {
8      int i = 1;
9      const long numIterations = 1000000;
10     std::thread t1{[&] {
11         for (int j = 0; j < numIterations; ++j) {
12             i++;
13         }
14     }};
15     std::thread t2{[&] {
16         for (int j = 0; j < numIterations; ++j) {
17             i--;
18         }
19     }};
20     t1.join();
21     t2.join();
22     std::cout << i << std::endl;
23 }
```

threadLamdaRace.cpp Output

```
1 $ ./threadLamdaRace
2 847046
3 $ ./threadLamdaRace
4 -18494
5 $ ./threadLamdaRace
6 25156
7 $ ./threadLamdaRace
8 -433633
9 $ ./threadLamdaRace
10 -8622
11 $ ./threadLamdaRace
12 -365615
13 $ ./threadLamdaRace
14 364766
15 $ ./threadLamdaRace
16 1
17 $ ./threadLamdaRace
18 1
19 $ ./threadLamdaRace
20 1000001
```

Atomic Operations

- Atomic types automatically synchronise and prevent race conditions
- The C++11 standard provides atomic primitive types:

Typedef name	Full specialisation
atomic_bool	atomic<bool>
atomic_char	atomic<char>
atomic_int	atomic<int>
atomic_uint	atomic<unsigned int>
atomic_long	atomic<long>

Atomic Operations Example

Our problematic race condition fixed with an atomic int:

```
1  #include <iostream>
2  #include <thread>
3  #include <atomic>
4  int main() {
5      std::atomic<int> i {1};
6      const long numIterations = 1000000;
7      std::thread t1([&i, numIterations] {
8          for (int j = 0; j < numIterations; ++j) {
9              i++;
10         }
11     });
12     std::thread t2([&i, numIterations] {
13         for (int j = 0; j < numIterations; ++j) {
14             i--;
15         }
16     });
17     t1.join();
18     t2.join();
19     std::cout << i << std::endl;
20 }
```

Is this a good approach?

Atomic Performance Problem

- The purpose of multi-threaded code is to improve performance
- Atomic operations slow down our code as they need to synchronise through locking
- How you structure multi-threaded code has a large influence on performance
- In our example we lock every time we need to increment, how might we do this better?

Improved Performance Example

```
1 int main() {
2     std::atomic<int> i {1};
3     const long numIterations = 1000000;
4     std::thread t1([&i,numIterations] {
5         int k = 0; // use a local variable
6         for (int j = 0; j < numIterations; ++j) {
7             k++; // generally more complex math
8         }
9         i += k;
10    });
11    std::thread t2([&i,numIterations] {
12        int k = 0;
13        for (int j = 0; j < numIterations; ++j) {
14            k--;
15        }
16        i -= k;
17    });
18    t1.join();
19    t2.join();
20    std::cout << i << std::endl;
21 }
```

Only write to the atomic int once per thread!

Atomic Operations Library

- The standard defines a number of operations on atomic types.
- These operations will be faster than your code.
- Example, `fetch_add()`:

```
1 #include <iostream>
2 #include <atomic>
3 #include <thread>
4
5 int main() {
6     std::atomic<int> i{10};
7     std::cout << "Initial Value: " << i << std::endl;
8     int fetchedValue = i.fetch_add(5);
9     std::cout << "Fetched Value: " << fetchedValue << std::endl;
10    std::cout << "Current Value: " << i << std::endl;
11 }
```

- Output:

```
1 Initial Value: 10
2 Fetched Value: 10
3 Current Value: 15
```

Atomic Classes

- The `std::atomic` library does not provide much support for creating atomic wrappers around your own data types.
- It can wrap a `TriviallyCopyable` type `T` which is a class with no user defined copy controllers.
- It can also wrap raw pointer types (but dereferences are not atomic!).
- See:
<http://en.cppreference.com/w/cpp/atomic/atomic>
- See: <http://stackoverflow.com/questions/15886308/stdatomic-with-custom-class-c-11>
- **Problem:** Synchronising on a whole class doesn't seem very efficient, what if only part of a class had race conditions?

C++11 Mutexes

- C++11 provides Mutex objects in the `<mutex>` header file.
- General idea:
 - A thread wants to read/write shared memory tries to lock the mutex object.
 - If another thread is currently locking the same mutex the first thread waits until the thread is unlocked or a timer expires.
 - When the thread obtains the lock it can safely read/write the shared memory
 - When the thread has finished using the shared memory it releases the lock
- Note:
 - If two or more threads are waiting for a lock to be released there is no order to which one will obtain the lock next.
 - It is assumed that all threads use the mutex to correctly access the shared memory.

std::mutex

- Non-timed mutex class
- Member functions:
 - `lock()` Tries to obtain the lock on the mutex and blocks indefinitely until the lock has been acquired.
 - `try_lock()` Tries to obtain the lock on the mutex, if the mutex is already locked will immediately return false, if the lock is obtained will return true.
 - `unlock()` Releases the lock currently held.
- **Note:** use `std::recursive_mutex` if your code is recursive and the same thread tries to repeatedly lock the same mutex.

std::mutex example

```
1 #include <iostream>
2 #include <thread>
3 #include <mutex>
4
5 int main() {
6     int i = 1;
7     const long numIterations = 1000000;
8     std::mutex iMutex;
9     std::thread t1([&] {
10         for (int j = 0; j < numIterations; ++j) {
11             iMutex.lock();
12             i++;
13             iMutex.unlock();
14         }
15     });
16     std::thread t2([&] {
17         for (int j = 0; j < numIterations; ++j) {
18             iMutex.lock();
19             i--;
20             iMutex.unlock();
21         }
22     });
23     t1.join();
24     t2.join();
25     std::cout << i << std::endl;
26 }
```

Mutex Performance Problem

- We now have the same problem we had with atomic operations!
- `std::mutex` locking takes exclusive ownership over objects.
- Our program slows down while threads are “busy waiting” to acquire mutex locks.
- Exclusive ownership of a lock over a critical section is required for safely writing data, but what about reading it?
- **Readers-writers locks** allow multiple threads to read data in parallel (shared ownership) but if data needs to be written exclusive ownership over the mutex must be acquired.

Readers-writers lock (C++14)

- C++14 `std::shared_timed_mutex` is a form of Readers-writers lock
- Shared ownership (read lock): `lock_shared()` blocks until the read lock is acquired. Multiple threads can `lock_shared()` the mutex in parallel.
- Exclusive ownership (write lock): `lock()` blocks until the write lock is acquired. Both the read and write locks need to be unlocked before ownership can be acquired.
- While a thread has exclusive ownership no other thread can obtain a lock (read or write).
- Unlocking: make sure you unlock the correct lock: `unlock()` and `unlock_shared()`

Exceptions and Locking

What happens to the mutex lock if an exception is thrown out of a critical section?

```
1  #include <iostream>
2  #include <thread>
3  #include <mutex>
4
5  int main() {
6      int i = 1;
7      const long numIterations = 1000000;
8      std::mutex iMutex;
9      std::thread t1([&] {
10         for (int j = 0; j < numIterations; ++j) {
11             iMutex.lock();
12             i++;
13             iMutex.unlock();
14         }
15     });
16     std::thread t2([&] {
17         for (int j = 0; j < numIterations; ++j) {
18             iMutex.lock();
19             if ( i == 0 ) throw ``i cannot be less than 0'';
20             i--;
21             iMutex.unlock();
22         }
23     });
24     t1.join();
25     t2.join();
26     std::cout << i << std::endl;
27 }
```


Exceptions and Locking

- **Major Problem:** the mutex lock is never released if we throw out of a critical section.
- Stack unwinding occurs until the exception is caught. But only inside this thread. If exception isn't caught then program terminates.
- Even if the exception is caught, if the mutex isn't unlocked then deadlocks and/or undefined program behaviour may occur.
- We've seen these types of issues before with dynamic memory.
How did we fix them?

Lock Guards

RAII wrapper class around a mutex.

```
1  #include <iostream>
2  #include <thread>
3  #include <mutex>
4
5  int main() {
6      int i = 1;
7      const long numIterations = 1000000;
8      std::mutex iMutex;
9      std::thread t1([&] {
10         for (int j = 0; j < numIterations; ++j) {
11             std::lock_guard<std::mutex> guard(iMutex);
12             i++;
13         }
14     });
15     std::thread t2([&] {
16         for (int j = 0; j < numIterations; ++j) {
17             std::lock_guard<std::mutex> guard(iMutex);
18             i--;
19         }
20     });
21     t1.join();
22     t2.join();
23     std::cout << i << std::endl;
24 }
```

Lock Guards

- Lock guards are similar to smart pointers.
- On construction they take ownership over (lock) a mutex.
- When they go out of scope they get destroyed. Their destructor unlocks the mutex.
- Are exception safe!
- You should never manually lock/unlock a mutex - always use lock guards

Lock guard limitations

- **Scenario** consider a box holding x elements:

```
1 struct Box {  
2     explicit Box(int num) : num_things{num} {}  
3  
4     int num_things;  
5     std::mutex m;  
6 };
```

- and a function to transfer items between two boxes.

```
1 void transfer(Box &from, Box &to, int num) {  
2     std::lock_guard<std::mutex> lock1(from.m);  
3     std::lock_guard<std::mutex> lock2(to.m);  
4  
5     from.num_things -= num;  
6     to.num_things += num;  
7 }
```

- **Problem** what happens if two threads in parallel attempt to move between the same two boxes in opposite directions?

Lock guard limitations

Problem: what happens if two threads in parallel attempt to move between the same two boxes in opposite directions?

```
1 int main() {  
2     Box acc1(100);  
3     Box acc2(50);  
4  
5     std::thread t1(transfer, std::ref(acc1), std::ref(acc2), 10);  
6     std::thread t2(transfer, std::ref(acc2), std::ref(acc1), 5);  
7  
8     t1.join();  
9     t2.join();  
10  
11     std::cout << acc1.num_things << std::endl;  
12 }
```

Remember: a `std::lock_guard` locks a mutex on construction and releases on destruction.

std::unique_lock

- Possible deadlocks can occur with multiple std::lock_guard objects
- std::unique_lock is a more complex lock guard
- Can defer locking to avoid deadlocks

```
1 void transfer(Box &from, Box &to, int num) {  
2     // don't actually lock the mutexs yet  
3     std::unique_lock<std::mutex> lock1(from.m, std::defer_lock);  
4     std::unique_lock<std::mutex> lock2(to.m, std::defer_lock);  
5  
6     // lock both unique_locks without deadlock  
7     std::lock(lock1, lock2);  
8  
9     from.num_things -= num;  
10    to.num_things += num;  
11 }
```

Example (modified) from:

http://en.cppreference.com/w/cpp/thread/unique_lock