

COMP6771

Advanced C++ Programming

Week 10

Multithreading - Producer/Consumer Problem

2016

www.cse.unsw.edu.au/~cs6771

Recap: C++11 Mutexes

- C++11 provides Mutex objects in the `<mutex>` header file.
- General idea:
 - A thread wants to read/write shared memory tries to lock the mutex object.
 - If another thread is currently locking the same mutex the first thread waits until the thread is unlocked or a timer expires.
 - When the thread obtains the lock it can safely read/write the shared memory
 - When the thread has finished using the shared memory it releases the lock

std::mutex

- Non-timed mutex class
- Member functions:
 - `lock()` Tries to obtain the lock on the mutex and blocks indefinitely until the lock has been acquired.
 - `try_lock()` Tries to obtain the lock on the mutex, if the mutex is already locked will immediately return false, if the lock is obtained will return true.
 - `unlock()` Releases the lock currently held.

std::mutex example

```
1 #include <iostream>
2 #include <thread>
3 #include <mutex>
4
5 int main() {
6     int i = 1;
7     const long numIterations = 1000000;
8     std::mutex iMutex;
9     std::thread t1([&] {
10         for (int j = 0; j < numIterations; ++j) {
11             iMutex.lock();
12             i++;
13             iMutex.unlock();
14         }
15     });
16     std::thread t2([&] {
17         for (int j = 0; j < numIterations; ++j) {
18             iMutex.lock();
19             i--;
20             iMutex.unlock();
21         }
22     });
23     t1.join();
24     t2.join();
25     std::cout << i << std::endl;
26 }
```

Lock Guards

RAII wrapper class around a mutex.

```
1 #include <iostream>
2 #include <thread>
3 #include <mutex>
4
5 int main() {
6     int i = 1;
7     const long numIterations = 1000000;
8     std::mutex iMutex;
9     std::thread t1([&] {
10         for (int j = 0; j < numIterations; ++j) {
11             std::lock_guard<std::mutex> guard(iMutex);
12             i++;
13         }
14     });
15     std::thread t2([&] {
16         for (int j = 0; j < numIterations; ++j) {
17             std::lock_guard<std::mutex> guard(iMutex);
18             i--;
19         }
20     });
21     t1.join();
22     t2.join();
23     std::cout << i << std::endl;
24 }
```

Producer-Consumer Problem

- **Scenario:** Consider a manufacturing plant which consumes the raw materials to produce products
- The manufacturing plant has limited storage space for both raw materials and final products.
- Trucks deliver raw materials to the plant, however, they arrive at random time intervals.
- If there is no space for the raw materials the trucks have to wait for space to become available.
- Trains remove the final products from the plant, if there are no products available they wait for one to be manufactured.

Producer-Consumer Problem

```
1 class ManufacturingPlant {  
2 public:  
3     ManufacturingPlant() : materialsCount{0} {}  
4  
5     void receiveMaterials(int i);  
6     int produceProduct();  
7 private:  
8     int materialsCount;  
9     std::mutex materialsCountMutex;  
10    const int CAPACITY = 100;  
11 };
```

Producer-Consumer Problem

```
1 class Truck {
2 public:
3     Truck(ManufacturingPlant& m) : mp{m} {}
4     void deliverMaterials() {
5         mp.receiveMaterials(10);
6     }
7 private:
8     ManufacturingPlant& mp;
9 };
10
11 class Train {
12 public:
13     Train(ManufacturingPlant& m) : mp{m} {}
14     void getProduct() {
15         mp.produceProduct();
16     }
17 private:
18     ManufacturingPlant& mp;
19 };
```


Producer-Consumer Problem

```
1 void ManufacturingPlant::receiveMaterials(int i) {
2     std::lock_guard<std::mutex> lg(materialsCountMutex);
3     if (materialsCount + i > CAPACITY) {
4         // TODO: wait for capacity!
5     }
6     materialsCount += i;
7 }
8
9 int ManufacturingPlant::produceProduct() {
10     std::lock_guard<std::mutex> lg(materialsCountMutex);
11     if (materialsCount - 10 > 0) {
12         // TODO: wait for materials to arrive
13     }
14     materialsCount -= 10;
15     return 10;
16 }
```

Problem: how do we release the lock and wait for either capacity or materials?

Condition Variables

- Condition variables allow threads to block, release a mutex, and wait until data is set by another thread or until a time period has elapsed.
- Explicit inter-thread communication.
- When we need to check/wait for a condition to be true, we call `wait()` on the condition variable.
- If we need to signal (communicate) that some data has been made available we need to `notify_one()` or `notify_all()` on the condition variable.
- Need to add to the class:

```
1 // used to signal the arrival of materials
2 std::condition_variable hasMaterials;
3 // used to signal the removal of materials
4 std::condition_variable hasCapacity;
```

Producer-Consumer with Condition Variables

```
1 void ManufacturingPlant::receiveMaterials(int i) {
2     std::unique_lock<std::mutex> lg(materialsCountMutex);
3     hasCapacity.wait(lg, [this, &i] {
4         if (materialsCount + i > CAPACITY) return false;
5         return true;
6     });
7     materialsCount += i;
8     hasMaterials.notify_one();
9 }
10
11 int ManufacturingPlant::produceProduct() {
12     std::unique_lock<std::mutex> lg(materialsCountMutex);
13     hasMaterials.wait(lg, [this] {
14         if (materialsCount - 10 < 0) return false;
15         return true;
16     });
17     materialsCount -= 10;
18     hasCapacity.notify_one();
19     return 10;
20 }
```

A further example is here:

<http://baptiste-wicht.com/posts/2012/04/>

Futures

- **Problem:** we've looked at lots of solutions to prevent deadlocks and memory corruption. But we're still stuck with busy waiting blocks.
- How do we make our code asynchronous and minimise waiting?
- e.g., If our train finds that it has to wait for a product to be manufactured, could it go and collect another product from a different manufacturing plant and then return to the first at a later period?

`std::async`

- `std::async` is used to create a thread and return a `std::future` which the result of the thread is stored in (e.g. it can work entirely entirely without shared memory and mutexes).
- `std::async` can hand over control of when to start a thread to the runtime system. It may not be called immediately if the system is busy.
- The `std::future` object can be used to check if the result of the `std::async` is available.

Example sketch

```
1  #include <iostream>
2  #include <future>
3  #include <thread>
4  #include <chrono>
5
6  int calculate() {
7      return 123;
8  }
9
10 int main() {
11     std::future<int> fut = std::async(calculate);
12     // note can force to launch a new thread using:
13     // std::future<int> fut = std::async(std::launch::async, calculate);
14
15     bool doOtherWork = true;
16     while (doOtherWork) {
17         // check if the result is available.
18         if (fut.wait_for(std::chrono::seconds(0)) == std::future_status::timeout) {
19             // do other work.. e.g. go to a different factory.
20         } else {
21             // either the result is available or the launch has been deferred
22             doOtherWork = false;
23         }
24     }
25
26     int res = fut.get(); // get the result from the future
27     std::cout << res << std::endl;
28 }
```

Modified example from: [http:](http://en.cppreference.com/w/cpp/thread/future/wait_for)

[//en.cppreference.com/w/cpp/thread/future/wait_for](http://en.cppreference.com/w/cpp/thread/future/wait_for)

Throwing exceptions across threads

- Futures can be used to transport exceptions across threads.
- When you call `get()` on the future you may get the result or the exception thrown.

```
1 int calculate() {  
2     throw std::runtime_error("Exception thrown from thread");  
3 }
```

```
1 try {  
2     int res = fut.get(); // get the result from the future  
3     std::cout << res << std::endl;  
4 } catch (const std::exception& ex) {  
5     std::cout << "Exception caught" << std::endl;  
6 }
```

Readings

- Chapter 23 Professional C++
- <http://baptiste-wicht.com/posts/2012/03/cp11-concurrency-tutorial-part-2-protect-shared-data.html>